

# FCS 150 - Project 4

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

# Goal

The goal of this project is to implement the support for a very simple *file system*: **ECS150-FS**.

Applications will have the possibility to read/write files from/to this file system.

```
int fs_mount(const char *diskname);           /* Mounting the file system */
int fs_umount(void);
int fs_info(void);

int fs_create(const char *filename);          /* Creating files */
int fs_delete(const char *filename);
int fs_ls(void);

int fs_open(const char *filename);            /* Opening files */
int fs_close(int fd);
int fs_stat(int fd);
int fs_lseek(int fd, size_t offset);

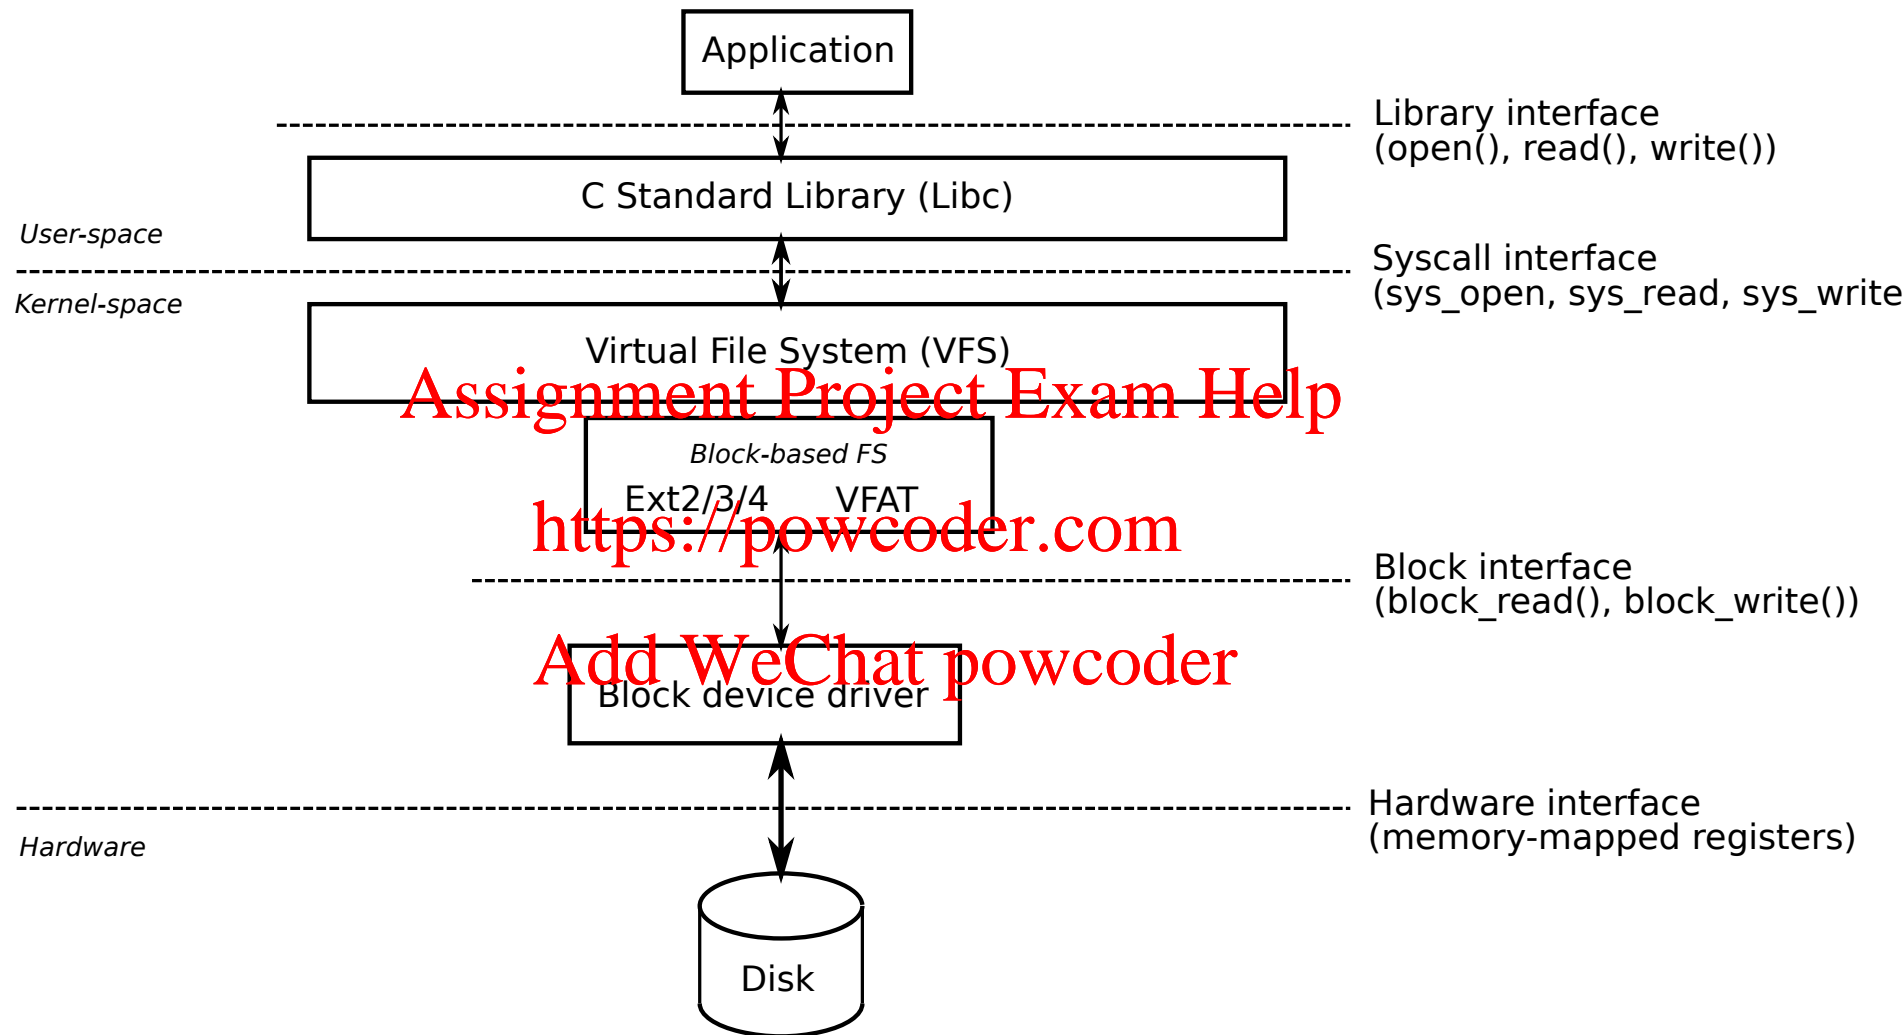
int fs_write(int fd, void *buf, size_t count); /* Modifying files */
int fs_read(int fd, void *buf, size_t count);
```

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

# Big picture: reality



Problem: the vast majority the file system management is in kernel mode!

# Emulating a disk with a file

A disk, or a partition on a disk, merely represents contiguous binary data storage.

How can we easily emulate any size of contiguous data?... With a file!

```
$ dd if=/dev/zero of=emulated_disk_space bs=4K count=8192
```

```
$ ls -l emulated_disk_space
```

```
-rw-r--r-- 1 joel joel 32M 2017-03-01 13:52 emulated_disk_space
```

```
$ xxd emulated_disk_space
```

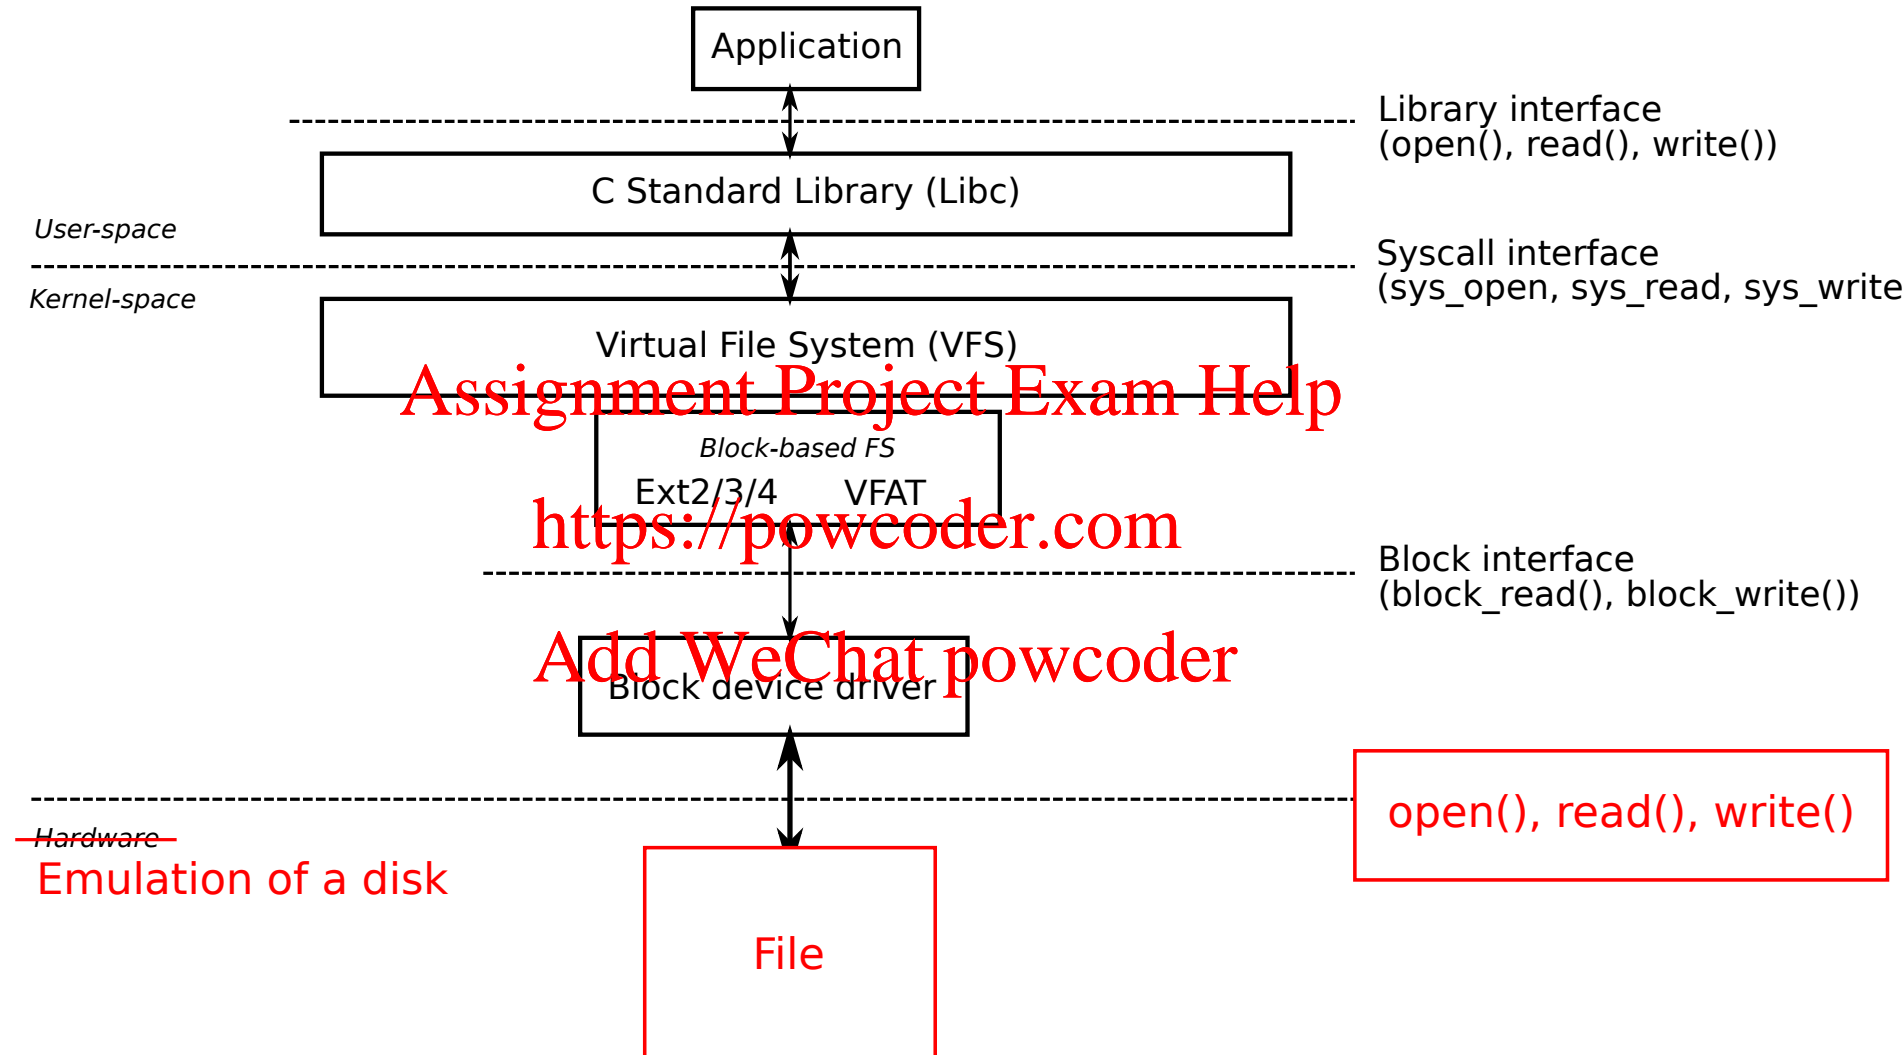
```
00000000: 0000 0000 0000 0000 0000 0000 0000 0000 .....
00000010: 0000 0000 0000 0000 0000 0000 0000 0000 .....
00000020: 0000 0000 0000 0000 0000 0000 0000 0000 .....
00000030: 0000 0000 0000 0000 0000 0000 0000 0000 .....
00000040: 0000 0000 0000 0000 0000 0000 0000 0000 .....
00000050: 0000 0000 0000 0000 0000 0000 0000 0000 .....
00000060: 0000 0000 0000 0000 0000 0000 0000 0000 .....
00000070: 0000 0000 0000 0000 0000 0000 0000 0000 .....
00000080: 0000 0000 0000 0000 0000 0000 0000 0000 .....
00000090: 0000 0000 0000 0000 0000 0000 0000 0000 .....
...
01ffffff: 0000 0000 0000 0000 0000 0000 0000 0000 .....
```

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

# Big picture: replacing the disk



# Accessing a file by blocks

```
#define BLOCK_SIZE 4096
```

```
int fd;
```

```
int block_open(char *disk_filename)
{
    fd = open(disk_filename, O_RDWR);
}
```

```
int block_read(size_t block_nr, void *buf)
{
    lseek(fd, block_nr * BLOCK_SIZE);
    read(fd, buf, BLOCK_SIZE);
}
```

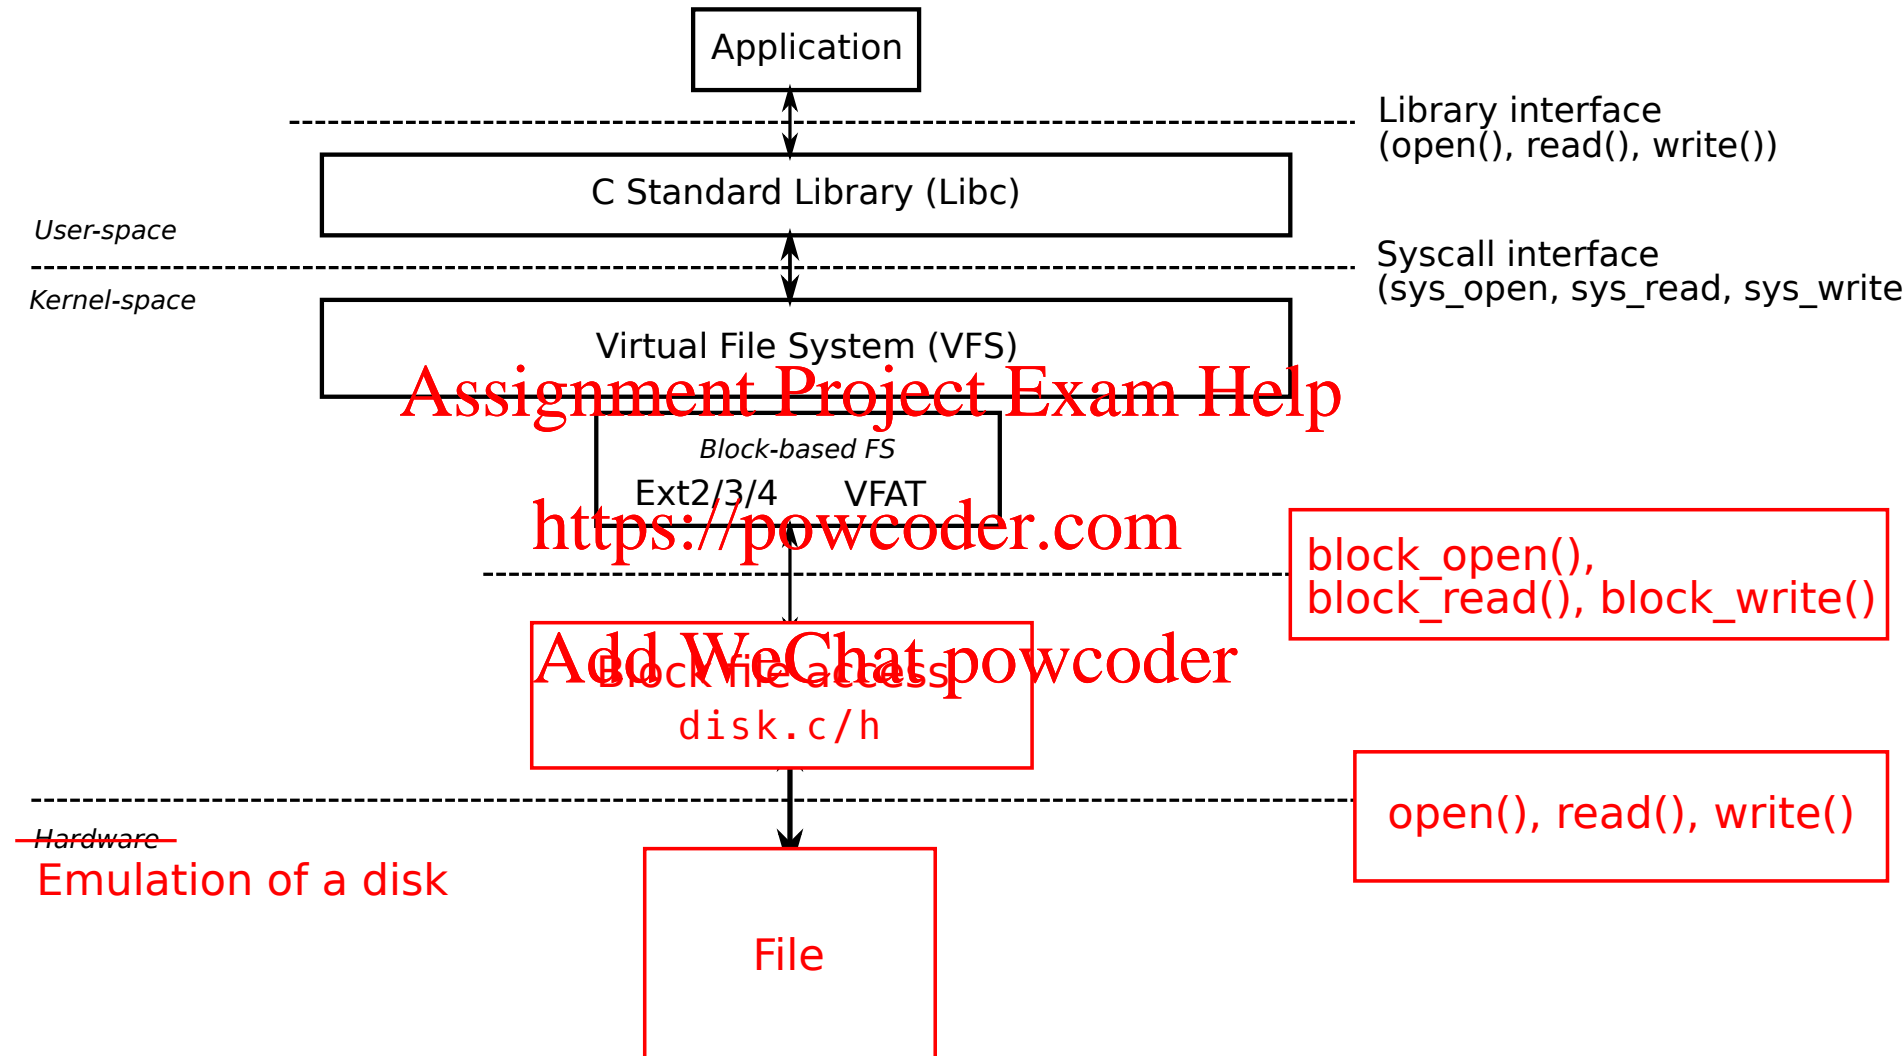
```
int block_write(size_t block_nr, void *buf)
{
    lseek(fd, block_nr * BLOCK_SIZE);
    write(fd, buf, BLOCK_SIZE);
}
```

Assignment Project Exam Help

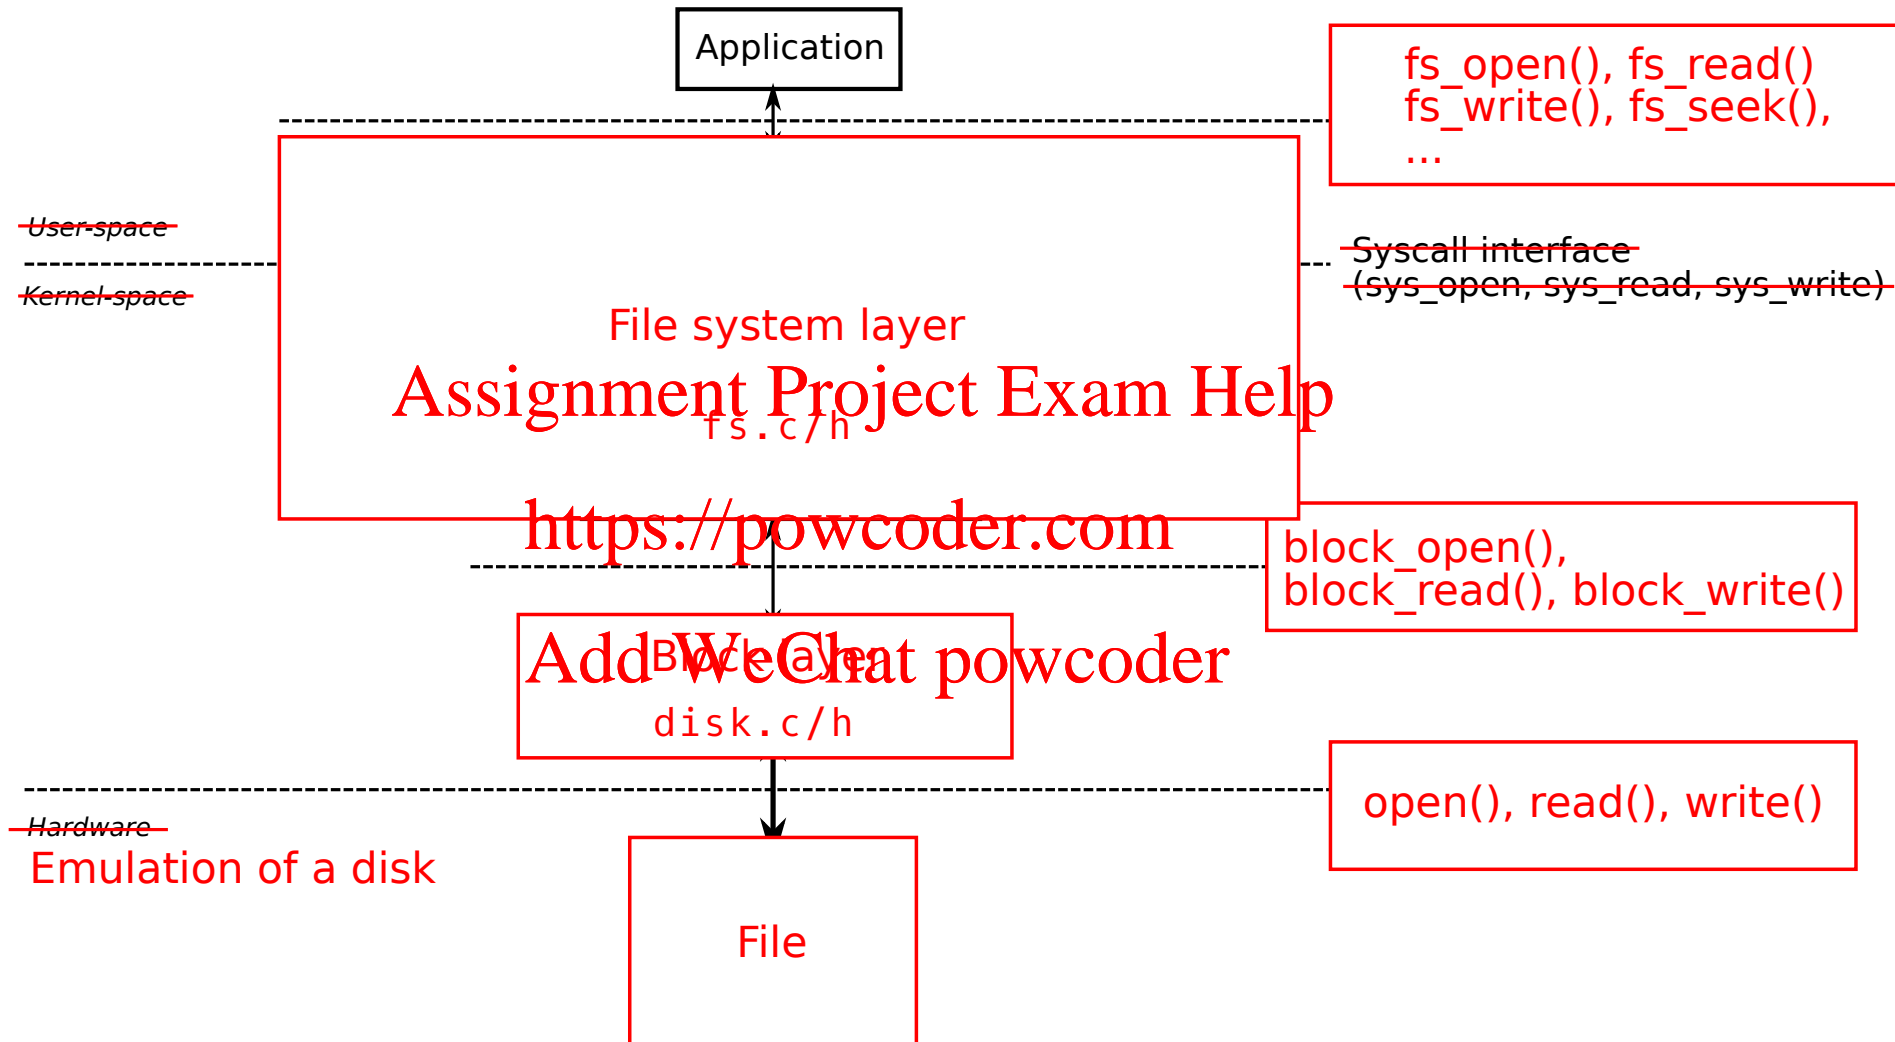
<https://powcoder.com>

Add WeChat powcoder

# Big picture: replacing the block device driver



# Big picture: replacing the libc/vfs/fs drivers





# ECS150-FS

## Layout



Each block is 4096 bytes.


Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

# ECS150-FS

## Layout



Each block is 4096 bytes.

Assignment Project Exam Help

Example with file system embedding 8192 data blocks:

- Amount of data blocks: 8192
- Number of blocks for FAT:  $(8192 * 2) / 4096 = 4$
- Total amount of blocks:  $1 + 4 + 1 + 8192 = 8198$
- Root directory block index: 5
- Data block start index: 6

<https://powcoder.com>

Add WeChat powcoder

# ECS150-FS

## Superblock: high-level layout

Offset	Length (bytes)	Description
0x00	8	Signature (must be equal to "ECS150FS")
0x08	2	Total amount of blocks of virtual disk
0x0A	2	Root directory block index
0x0C	2	Data block start index
0x0E	2	Amount of data blocks
0x10	1	Number of blocks for FAT
0x11	4079	Unused/Padding

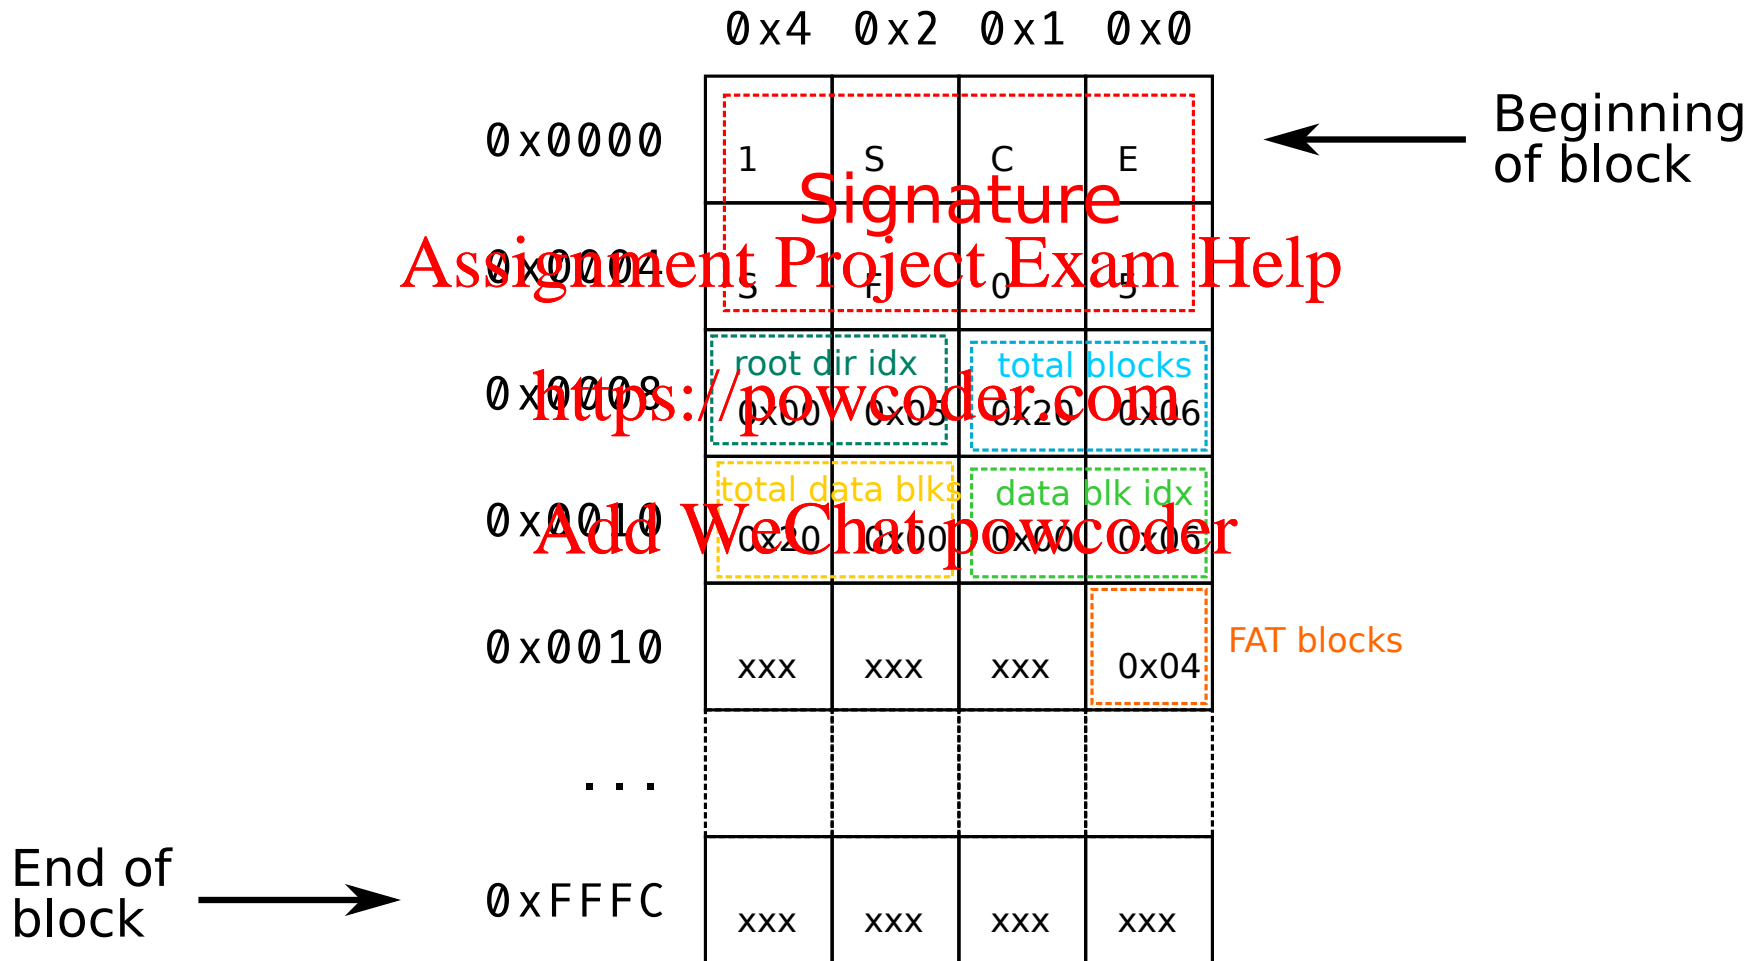
Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

# ECS150-FS

## Superblock: at byte level



# ECS150-FS

## Superblock: C data structure

```
struct superblock{  
    ???  
};
```

Key points:

- The integer types must match exactly those of the specification
- Careful about alignment

<https://powcoder.com>

Add WeChat powcoder

# Digression

## Integer types

- Is `char` always 8 bits?
- Is `short int` always 16 bits?
- Is `int` always 32 bits?
- Etc.

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

# Digression

## Integer types

- Is `char` always 8 bits?
- Is `short int` always 16 bits?
- Is `int` always 32 bits?
- Etc.

Type	Specification
char	"Smallest addressable unit of the machine that can contain basic character set"
short	"Capable of containing <i>at least</i> the $[-32767, +32767]$ range; thus, it is <i>at least</i> 16 bits in size."
int	"Capable of containing <i>at least</i> the $[-32767, +32767]$ range; thus, it is <i>at least</i> 16 bits in size."
long	"Capable of containing <i>at least</i> the $[-2147483647, +2147483647]$ range; thus, it is <i>at least</i> 32 bits in size."

How to guarantee a certain size then?

# Digression

## Integer types

Use integer types that have exact widths:

```
#include <stdint.h>
```

```
int8_t  
int16_t  
int32_t
```

```
uint8_t  
uint16_t  
uint32_t
```

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder



# Digression

## Structure alignment

Structure naturally packed:

```
struct packed_s  
{  
    int32_t    a;  
    int16_t    b;  
    int16_t    c;  
};
```

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

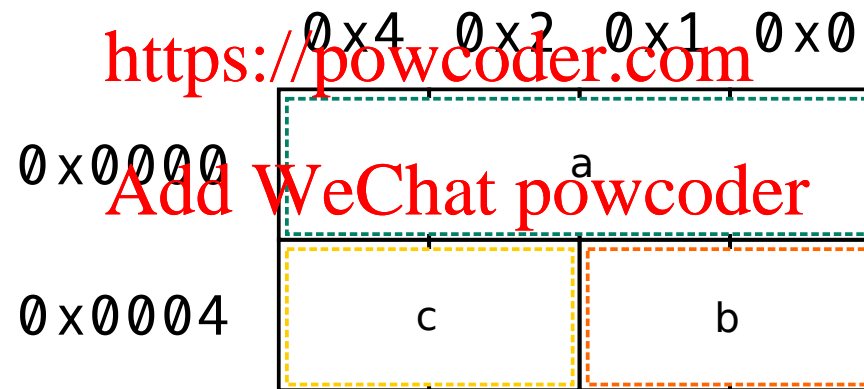
# Digression

## Structure alignment

Structure naturally packed:

```
struct packed_s
{
    int32_t    a;
    int16_t    b;
    int16_t    c;
};
```

Assignment Project Exam Help



# Digression

## Structure alignment

Structure fields have to aligned...

```
struct padded_s
{
    int8_t    a;
    int32_t   b;
    int16_t   c;
};
```

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

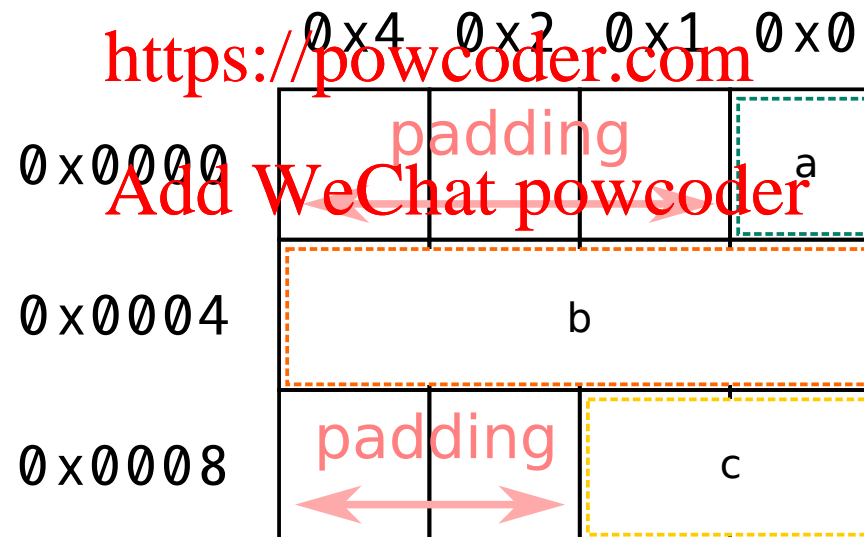
# Digression

## Structure alignment

Structure fields have to aligned...

```
struct padded_s
{
    int8_t    a;
    int32_t   b;
    int16_t   c;
};
```

Assignment Project Exam Help



# Digression

## Structure alignment

Force compiler to ignore alignment

```
struct packed_s  
{  
    int8_t    a;  
    int32_t   b;  
    int16_t   c;  
} __attribute__((packed));
```

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

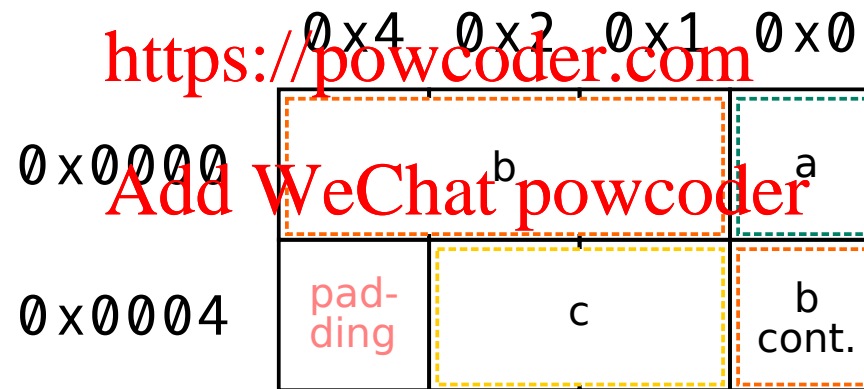
# Digression

## Structure alignment

Force compiler to ignore alignment

```
struct packed_s  
{  
    int8_t    a;  
    int32_t   b;  
    int16_t   c;  
} __attribute__((packed));
```

Assignment Project Exam Help



# Digression

## Structure alignment

Conclusion: when transposing a specification into data structures, always use packing!

- File format
- Network protocol
- Etc.

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

# Digression

## Reading data structures from a file (or buffer)

Reading data from a file (or whatever blob of data), for which I know the layout. It can easily be type-casted into a structure instance.

```
struct packed_s
{
    int32_t  a;
    int16_t  b;
    int16_t  c;
};
```

```
char* buf[8];

fd = open("file", O_RDWR);
read(bd, buf, 8);


struct packed_s *s = buf;
s->a = 0;

/* or simply */
struct packed_s obj;
read(bd, &obj, sizeof(obj));
obj.a = 0;
```

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder





# ECS150-FS

## FAT

- Big array of 16-bit entries: linked-list of data blocks composing a file
- Three possible values for each entry:
  - 0: corresponding data block is available
  - FAT\_EOC: last data block of a file
  - !=0 && !=FAT\_EOC: index of next data block

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

# ECS150-FS

## FAT

- Big array of 16-bit entries: linked-list of data blocks composing a file
- Three possible values for each entry:
  - 0: corresponding data block is available
  - FAT\_EOC: last data block of a file
  - !=0 && !=FAT\_EOC: index of next data block

## Root directory

1 block, 16-byte entry per file: 128 entries total

Offset	Length (bytes)	Description
0x00	16	Filename (including NULL character)
0x10	4	Size of the file (in bytes)
0x14	2	Index of the first data block
0x16	10	Unused/Padding

# ECS150-FS

Example: big file, small file, empty file

<test1, 25000, 2>,  
<test2, 5000, 1>,  
<test3, 0, FAT\_EOC>  
...

FAT

0	
1	8
2	3
3	4
4	5
5	6
6	7
7	0xFFFF
8	0xFFFF
9	0
10	0

Data blocks

0	
1	test2, block #0
2	test1, block #0
3	test1, block #1
4	test1, block #2
5	test1, block #3
6	test1, block #4
7	test1, block #5
8	test2, block #1
9	
10	

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

# Implementation

## Phase 1: Volume mounting

- `fs_mount()`: open the virtual disk, and read the metadata (superblock, fat, root directory)
- `fs_unmount()`: close virtual disk (make sure that virtual disk is up-to-date)
- `fs_info()`: show information about volume

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

# Implementation

## Phase 2: File creation/deletion

- `fs_create()`: Create a new file
  - Initially, size is 0 and pointer to first data block is `FAT_EOC`
- `fs_delete()`: Delete an existing file
  - Don't forget to free allocated data blocks
- `fs_ls()`: List all the existing files

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

# Implementation

## Phase 3: File descriptor operations

- `fs_open()`: initialize and return file descriptor
  - 32 file descriptors max
  - Can open same file multiple times
  - Contains file's offset (initially 0)
- `fs_close()`: close file descriptor
- `fs_seek()`: move file's offset
- `fs_stat()`: return file's size

None of these function should change the file system...

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

# Implementation

## Phase 4: File reading/writing

Most complicated phase: might take as much time as all the previous phases combined

- Allocation of new blocks must follow *first-fit* strategy (allocate first free data block from beginning of the FAT).
- Three difficulties:
  - Small operations
  - First/last block on big operations
  - Extending writes

Assignment Project Exam Help

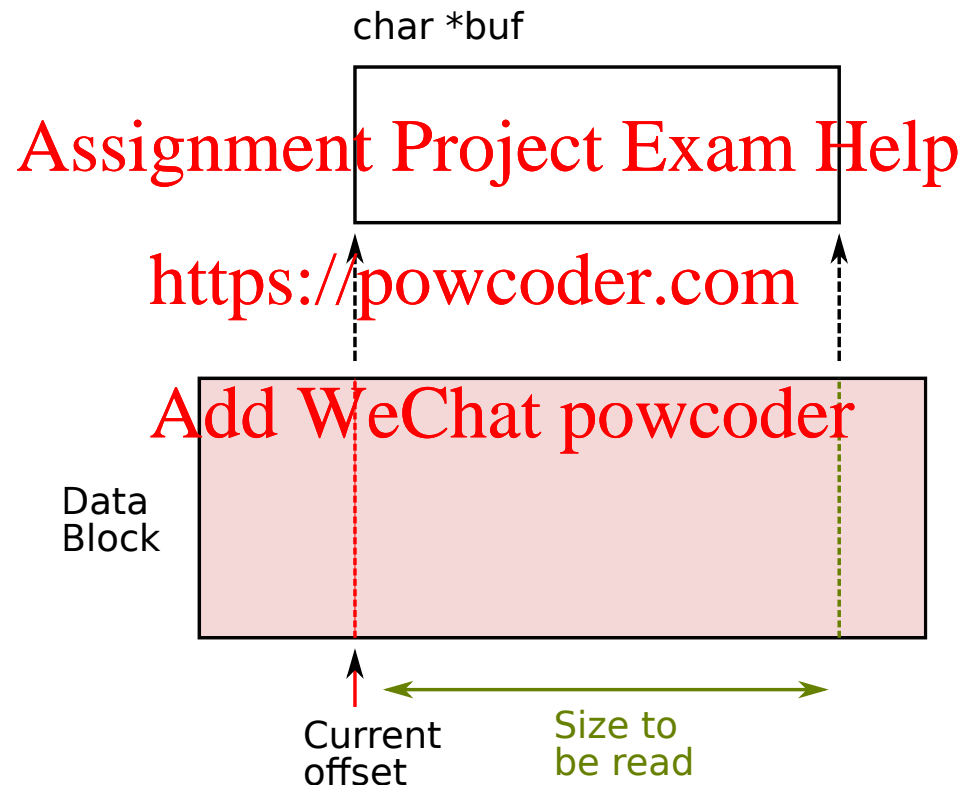
<https://powcoder.com>

Add WeChat powcoder

# Implementation

## Small operation: example

- Current offset is in the middle of the file, not aligned on the beginning of a block
- The size of data to read is smaller than what's remaining in this block



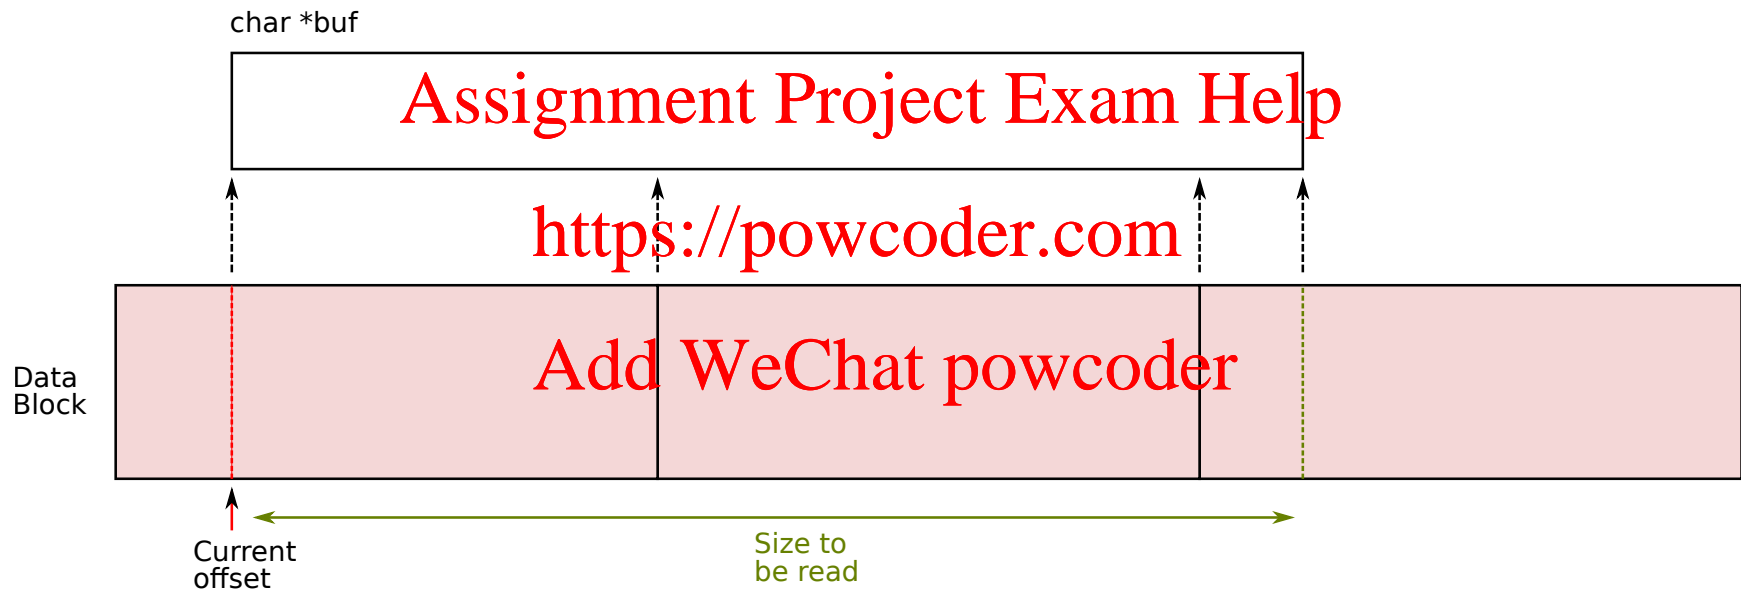
Might want to use a *bounce buffer*



# Implementation

## Big operation: example

- Current offset is in the middle of the file, not aligned on the beginning of a block
- The size to read spans multiple (non-consecutive) blocks
- The size of data to read is smaller than what's remaining in the last block



Mix of *bounce buffer* and direct copy

# Implementation

## Extending write: example

- Write more than what's currently allocated



# Implementation

## In short

- Think of all the cases: combination of file's offset, file's size, size to be read or written, etc.
- Come up with a way to handle all these combinations in the *most* generic way (ie not one function per case!)

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder