

ECS 150: Project #4 - File system

Dr. Joël Porquet

UC Davis, Spring Quarter 2017, version 0

- [Changelog](#)
- [General information](#)
- [Specifications](#)
 - [Introduction](#)
 - [Constraints](#)
 - [Skeleton code](#)
 - [Phase 0: the ECS150-FS specification](#)
 - [Superblock](#)
 - [FAT](#)
 - [Root directory](#)
 - [Formatting program](#)
 - [Reference program and testing](#)
 - [Phase 1: Mounting/unmounting](#)
 - [Phase 2: File creation/deletion](#)
 - [Phase 3: File descriptor operations](#)
 - [Phase 4: File reading/writing](#)
- [Deliverable](#)
 - [Content](#)
 - [Git](#)
 - [Handin](#)
- [Academic integrity](#)

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

Changelog

- v0: First publication

General information

Due before 11:59 PM, Friday, June 2nd, 2017.

You will be working with a partner for this project. Remember, you cannot keep the same partner for more than 2 projects over the course of the quarter.

The reference work environment is the CSIF.

Specifications

Note that the specifications for this project are subject to change at anytime for additional clarification. Make sure to always refer to the latest version.

Introduction

The goal of this project is to implement the support of a very simple file system, **ECS150-FS**. This file system is based on a FAT and supports up to 128 files in a single root directory.

The file system is implemented on top of a virtual disk. This virtual disk is actually a simple binary file that is stored on the "real" file system provided by your computer.

Exactly like real hard drives which are split into sectors, the virtual disk is logically split into blocks. The first software layer involved in the file system implementation is the *block API* and is provided to you. This block API is used to open or close a virtual disk, and read or write entire blocks from it.

Above the block layer, the *FS layer* is in charge of the actual file system management. Through the FS layer, you can mount a virtual disk, list the files that are part of the disk, add or delete new files, read from files or write to files, etc.

Constraints

Your library must be written in C, be compiled with GCC and only use the standard functions provided by the [GNU C Library](#) (aka `libc`). All the functions provided by the `libc` can be used, but your program cannot be linked to any other external libraries.

Your source code should follow the relevant parts of the [Linux kernel coding style](#) and be properly commented.

Assignment Project Exam Help

Skeleton code

The skeleton code that you are expected to complete is available in `/home/jporquet/ecs150/`. This code already defines most of the prototypes for the functions you must implement, as explained in the following sections.

```
$ cd /home/jporquet/ecs150/
$ tree
```

```
.
├── libfs
│   ├── disk.c
│   ├── disk.h
│   ├── fs.c*
│   ├── fs.h
│   └── Makefile*
├── Makefile
└── test_fs.c
```

In the subdirectory `libfs`, there are the files composing the file-system library that you must complete. The files to complete are marked with a star (you should have **no** reason to touch any of the headers which are not marked with a star, even if you think you do...).

Phase 0: the ECS150-FS specification

For this project, the specification for a very simple file system has been defined: it's the **ECS150-FS** file system.

The layout of ECS150-FS on a disk is composed of four consecutive logical parts:

- The *Superblock* is the very first block of the disk and contains information about the file system (number of blocks, size of the FAT, etc.)
- The *File Allocation Table* is located on one or more blocks, and keeps track of both the free data

- blocks and the mapping between files and the data blocks holding their content.
- The *Root directory* is in the following block and contains an entry for each file of the file system, defining its name, size and the location of the first data block for this file.
- Finally, all the remaining blocks are *Data blocks* and are used by the content of files.

The size of virtual disk blocks is **4096 bytes**.

Unless specified otherwise, all values in this specification are *unsigned* and the order of multi-bytes values is *little-endian*.

Superblock

The superblock is the first block of the file system. Its internal format is:

Offset	Length (bytes)	Description
0x00	8	Signature (must be equal to "ECS150FS")
0x08	2	Total amount of blocks of virtual disk
0x0A	2	Root directory block index
0x0C	2	Data block start index
0x0E	2	Amount of data blocks
0x10	1	Number of blocks for FAT
0x11	4079	Unused/Padding

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

If one creates a file system with 8192 data blocks, the size of the FAT will be $8192 \times 2 = 16384$ bytes long, thus spanning $16384 / 4096 = 4$ blocks. The root directory block index will therefore be 5, because before it there are the superblock (block index #0) and the FAT (starting at block index #1 and spanning 4 blocks). The data block start index will be 6, because it's located right after the root directory block. The total amount of blocks for such a file system would then be $1 + 4 + 1 + 8192 = 8198$.

FAT

The FAT is a flat array, possibly spanning several blocks, which entries are composed of 16-bit unsigned words. There are as many entries as *data blocks* in the disk.

The first entry of the FAT (entry #0) is always invalid and contains the special *FAT_EOC (End-of-Chain)* value which is $0xFFFF$. Entries marked as 0 correspond to free data blocks. Entries containing a positive value are part of a chainmap and represent a link to the next block in the chainmap.

Note that although the numbering in the FAT starts at 0, entry contents must be added to the data block start index in order to find the real block number on disk.

The following table shows an example of a FAT containing two files:

- The first file is of length 18,000 bytes (thus spanning 5 data blocks) and is contained in consecutive data blocks (DB#2, DB#3, DB#4, DB#5, DB#6 and DB#7).
- The second file is of length 5,000 bytes (this spanning 2 data blocks) and its content is fragmented in two non-consecutive data blocks (DB#1 and DB#8).

Each entry in the FAT is 16-bit wide.

FAT index: 0 1 2 3 4 5 6 7 8 9 10 ...

Content: 0xFFFF 8 3 4 5 6 7 0xFFFF 0xFFFF 0 0 ...

Root directory

The root directory is an array of 128 entries stored in the block following the FAT. Each entry is 32-byte wide and describes a file, according to the following format:

Offset	Length (bytes)	Description
0x00	16	Filename (including NULL character)
0x10	4	Size of the file (in bytes)
0x14	2	Index of the first data block
0x16	10	Unused/Padding

An empty entry is defined by the first character of the entry's filename being equal to 0.

The entry for an empty file, which doesn't have any data blocks, would have its size be 0, and the index of the first data block be `FAT_EOC`.

Continuing the previous example, let's assume that the first file is named "test1" and the second file "test2". Let's also assume that there is an empty file named "test3". The content of the root directory would be:

Filename (16 bytes)	Size (4 bytes)	Index (2 bytes)	Padding (10 bytes)
---------------------	----------------	-----------------	--------------------

test1	18000	2	xxx
test2	5000	1	xxx
test3	0	FAT_EOC	xxx
	xxx	xxx	xxx
...

Formatting program

A FS formatter is provided to you in `/home/jporquet/ecs150/fs_make.x`. The purpose of this program is to create a new virtual disk and initialize an empty file system on it.

It accepts two arguments: the name of the virtual disk image and the number of data blocks to create:

```
$ ./fs_make.x disk.fs 4096
Created virtual disk 'disk.fs' with '4096' data blocks
```

Note that this formatter can not create a file-system of more than 8192 data blocks (which already makes

a 33MB virtual disk).

Reference program and testing

A reference program is provided to you in `/home/jporquet/ecs150/fs_ref.x`. This program accepts multiple commands (one per run) and allows you to:

- Get some information about a virtual disk: `$ fs_ref.x info <diskname>`
- List all the files contained in a virtual disk: `$ fs_ref.x ls <diskname>`
- Etc. in order to have the list of commands: `$ fs_ref.x`

The code of this executable is actually provided to you in `test_fs.c` and you will have to implement the complete API that this program uses. The creation of new virtual disks is the only tasks that you don't have to program yourself as it is provided by `fs_make.x`.

Otherwise your implementation should generate the same output as the reference program, and the manipulations that are performed on the virtual disk should be understood by both the reference program and your implementation. For example, after creating a virtual disk, the output of the command `info` from the reference program should match exactly the output from your implementation:

```
$ ./fs_make.x disk.fs 8192
Creating virtual disk 'disk.fs' with '8192' data blocks
$ ./fs_ref.x info disk.fs > ref_output
$ ./test_fs.x info disk.fs > my_output
$ diff ref_output my_output
$
```

Assignment Project Exam Help

Phase 1: Mounting/unmounting

<https://powcoder.com>

In this first phase, you must implement the function `fs_mount()` and `fs_umount()`.

`fs_mount()` makes the file system contained in the specified virtual disk "ready to be used". You need to open the virtual disk, using the block API, and load the meta-information that is necessary to handle the file system operations described in the following phases. `fs_umount()` makes sure that the virtual disk is properly closed and that all the internal data structures of the FS layer are properly cleaned.

Add WeChat powcoder

For this phase, you should probably start by defining the data structures corresponding to the blocks containing the meta-information about the file system (superblock, FAT and root directory).

In order to correctly describe these data structures, you will probably need to use the integer types defined in `stdint.h`, such as `int8_t`, `uint8_t`, `uint16_t`, etc. Also, when describing your data structures and in order to avoid the compiler to interfere with their layout, it's always good practice to attach the attribute `packed` to these data structures.

Don't forget that your function `fs_mount()` should perform some error checking in order to verify that the file system has the expected format. For example, the signature of the file system should correspond to the one defined by the specification, the total amount of block should correspond to what `block_disk_count()` returns, etc.

Once you're able to mount a file system, you can implement the function `fs_info()` which prints some information about the mounted file system and make sure that the output corresponds exactly to the reference program.

It is important to observe that the file system must provide persistent storage. Let's assume that you have created a file system on a virtual disk and mounted it. Then, you create a few files and write some data to them. Finally, you unmount the file system. At this point, all data must be written onto the virtual disk.

Another application that mounts the file system at a later point in time must see the previously created files and the data that was written. This means that whenever `umount_fs()` is called, all meta-information and file data must have been written out to disk.

Phase 2: File creation/deletion

In this second phase, you must implement `fs_create()` and `fs_delete()` which add or remove files from the file system.

In order to add a file, you need to find an empty entry in the root directory and fill it out with the proper information. At first, you only need to specify the name of the file and reset the other information since there is no content at this point. The size should be set to 0 and the first index on the data blocks should be set to `FAT_EOC`.

Removing a file is the opposite procedure: the file's entry must be emptied and all the data blocks containing the file's contents must be freed in the FAT.

Once you're able to add and remove files, you can implement the function `fs_ls()` which prints the listing of all the files in the file system. Make sure that the output corresponds exactly to the reference program.

Phase 3: File descriptor operations

In order for applications to manipulate files, the FS API offers functions which are very similar to the Linux file system operations. `fs_open()` opens a file and returns a *file descriptor* which can then be used for subsequent operations on this file (reading, writing, changing the file offset, etc). `fs_close()` closes a file descriptor.

Your library must support a maximum of 32 file descriptors that can be open simultaneously. The same file (i.e. file with the same name) can be opened multiple times, in which case `fs_open()` must return multiple independent file descriptors.

A file descriptor is associated to a file and also contains a *file offset*. The file offset indicates the current reading/writing position in the file. It is implicitly incremented whenever a read or write is performed, or can be explicitly set by `fs_lseek()`.

Finally, the function `fs_stat()` must return the size of the file corresponding to the specified file descriptor. To append to a file, one can, for example, call `fs_lseek(fd, fs_stat(fd));`.

Phase 4: File reading/writing

In the last phase, you must implement `fs_read()` and `fs_write()` which respectively read from and write to a file.

It is advised to start with `fs_read()` which is slightly easier to program. You can use the reference program to write a file in a disk image, which you can then read using your implementation.

For these functions, you will probably need a few helper functions. For example, you will need a function that returns the index of the data block corresponding to the file's offset. For writing, in the case the file has to be extended in size, you will need a function that allocates a new data block and link it at the end of the file's data block chain. Note that the allocation of new blocks should follow the *first-fit* strategy (first block available from the beginning of the FAT).

When reading or writing a certain number of bytes from/to a file, you will also need to deal properly with

possible "mismatches" between the file's current offset, the amount of bytes to read/write, the size of blocks, etc.

For example, let's assume a reading operation for which the file's offset is not aligned to the beginning of the block or the amount of bytes to read doesn't span the whole block. You will probably need to read the entire block into a *bounce* buffer first, and then copy only the right amount of bytes from the bounce buffer into the user-supplied buffer.

The same scenario for a writing operation would be slightly trickier. You will probably need to first read the block from disk, then modify only the part starting from the file's offset with the user-supplied buffer, before you can finally write the dirty block back to the disk.

These special cases happen mostly for small reading/writing operations, or at the beginning or end of a big operation. In big operations (spanning multiple blocks), offsets and sizes are perfectly aligned for all the middle blocks and the procedure is then quite simple, as blocks can be read or written entirely.

Deliverable

Content

Your submission should contain, besides your source code, the following files:

- **AUTHORS:** first name, last name, student ID, CSIF username and email of each partner, one entry per line formatted in CSV (fields are separated with commas). For example:

```
$ cat AUTHORS
Homer,Simpson,00010001,simpson32,hsimpson@ucdavis.edu
Robert David,Martin,00010002,rdm,rdmartin@ucdavis.edu
```

- **REPORT.md:** a description of your submission. Your report must respect the following rules:
 - It must be formatted in markdown language as described in this [Markdown-Cheatsheet](#).
 - It should contain no more than 300 lines and the maximum width for each line should be 80 characters (check your editor's settings to configure that).
 - It should explain your high-level design choices, details about the relevant parts of your implementation, how you tested your project, the sources that you may have used to complete this project, and any other information that can help understanding your code.
 - Keep in mind that the goal of this report is not to paraphrase the assignment, but to explain how you implemented it.
- **libfs/Makefile:** a Makefile that compiles your source code without any errors or warnings (on the CSIF computers), and builds a static library named `libfs.a`.

The compiler should be run with the options `-Wall -Werror`.

There should also be a `clean` rule that removes generated files and puts the directory back in its original state.

Your submission should be empty of any clutter files such as executable files, core dumps, etc.

Git

Your submission must be under the shape of a Git bundle. In your git repository, type in the following command (your work must be in the branch `master`):

```
$ git bundle create fs.bundle master
```

It should create the file `fs.bundle` that you will submit via `handin`.

You can make sure that your bundle has properly been packaged by extracting it in another directory and verifying the log:

```
$ cd /path/to/tmp/dir
$ git clone /path/to/fs.bundle -b master fs
$ cd fs
$ ls -l
...
$ git log
...
```

Handin

Your Git bundle, as created above, is to be submitted with `handin` from one of the CSIF computers:

```
$ handin cs150 p4 fs.bundle
Submitting fs.bundle... ok
$
```

You can verify that the bundle has been properly submitted:

```
$ handin cs150 p4
The following input files have been received:
...
$
```

Academic integrity

You are expected to write this project from scratch, thus avoiding to use any existing source code available on the Internet. You must specify in your `REPORT.md` file any sources of code that you or your partner have viewed to help you complete this project. All class projects will be submitted to MOSS to determine if pairs of students have excessively collaborated with other pairs, or have used the work of past students.

Excessive collaboration, or failure to list external code sources will result in the matter being transferred to Student Judicial Affairs.